

# Game playing

## CHAPTER 5, SECTIONS 1–5

# Outline

- ◇ Perfect play
- ◇ Resource limits
- ◇  $\alpha$ - $\beta$  pruning
- ◇ Games of chance

## Games vs. search problems

“Unpredictable” opponent  $\Rightarrow$  solution is a contingency plan

Time limits  $\Rightarrow$  unlikely to find goal, must approximate

Plan of attack:

- algorithm for perfect play (Von Neumann, 1944)
- finite horizon, approximate evaluation (Zuse, 1945; Shannon, 1950; Samuel, 1952–57)
- pruning to reduce costs (McCarthy, 1956)

## Types of games


	deterministic	chance
perfect information	chess, checkers, go, othello	backgammon monopoly
imperfect information		bridge, poker, scrabble nuclear war

# Minimax

Perfect play for deterministic, perfect-information games

Idea: choose move to position with highest *minimax value*  
= best achievable payoff against best play

E.g., 2-ply game:



# Minimax algorithm

**function** MINIMAX-DECISION(*game*) *returns an operator*

**for each** *op* **in** OPERATORS[*game*] **do**

    VALUE[*op*] ← MINIMAX-VALUE(APPLY(*op*, *game*), *game*)

**end**

**return** the *op* with the highest VALUE[*op*]

---

**function** MINIMAX-VALUE(*state*, *game*) *returns a utility value*

**if** TERMINAL-TEST[*game*](*state*) **then**

**return** UTILITY[*game*](*state*)

**else if** MAX is to move in *state* **then**

**return** the highest MINIMAX-VALUE of SUCCESSORS(*state*)

**else**

**return** the lowest MINIMAX-VALUE of SUCCESSORS(*state*)

# Properties of minimax

Complete??

Optimal??

Time complexity??

Space complexity??

## Properties of minimax

Complete?? Yes, if tree is finite (chess has specific rules for this)

Optimal?? Yes, against an optimal opponent. Otherwise??

Time complexity??  $O(b^m)$

Space complexity??  $O(bm)$  (depth-first exploration)

For chess,  $b \approx 35$ ,  $m \approx 100$  for “reasonable” games  
 $\Rightarrow$  exact solution completely infeasible




## Resource limits

Suppose we have 100 seconds, explore  $10^4$  nodes/second  
 $\Rightarrow$   $10^6$  nodes per move

Standard approach:


- *cutoff test*  
e.g., depth limit (perhaps add *quiescence search*)
- *evaluation function*  
= estimated desirability of position

# Evaluation functions



Black to move

White slightly better



White to move

Black winning

For chess, typically *linear* weighted sum of features


$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

e.g.,  $w_1 = 9$  with

$f_1(s) = (\text{number of white queens}) - (\text{number of black queens})$

etc.

## Digression: Exact values don't matter



Behaviour is preserved under any *monotonic* transformation of  $E_{VAL}$

Only the order matters:

payoff in deterministic games acts as an *ordinal utility* function

## Cutting off search

MINIMAXCUTOFF is identical to MINIMAXVALUE except

1. TERMINAL? is replaced by CUTOFF?
2. UTILITY is replaced by EVAL

Does it work in practice?

$$b^m = 10^6, \quad b = 35 \quad \Rightarrow \quad m = 4$$


4-ply lookahead is a hopeless chess player!


4-ply  $\approx$  human novice


8-ply  $\approx$  typical PC, human master


12-ply  $\approx$  Deep Blue, Kasparov

## $\alpha$ - $\beta$ pruning example












**X 3**

**X 2**



## Properties of $\alpha-\beta$

Pruning *does not* affect final result

Good move ordering improves effectiveness of pruning


With “perfect ordering,” time complexity =  $O(b^{m/2})$

⇒ *doubles* depth of search

⇒ can easily reach depth 8 and play good chess

A simple example of the value of reasoning about which computations are relevant (a form of *metareasoning*)

## Why is it called $\alpha$ - $\beta$ ?



$\alpha$  is the best value (to MAX) found so far off the current path

If  $V$  is worse than  $\alpha$ , MAX will avoid it  $\Rightarrow$  prune that branch

Define  $\beta$  similarly for MIN

# The $\alpha$ - $\beta$ algorithm

Basically MINIMAX + keep track of  $\alpha$ ,  $\beta$  + prune

**function** MAX-VALUE( $state, game, \alpha, \beta$ ) **returns** the minimax value of  $state$

**inputs:**  $state$ , current state in game

$game$ , game description

$\alpha$ , the best score for MAX along the path to  $state$

$\beta$ , the best score for MIN along the path to  $state$

**if** CUTOFF-TEST( $state$ ) **then return** EVAL( $state$ )

**for each**  $s$  **in** SUCCESSORS( $state$ ) **do**

$\alpha \leftarrow \text{MAX}(\alpha, \text{MIN-VALUE}(s, game, \alpha, \beta))$

**if**  $\alpha \geq \beta$  **then return**  $\beta$

**end**

**return**  $\alpha$

---

**function** MIN-VALUE( $state, game, \alpha, \beta$ ) **returns** the minimax value of  $state$

**if** CUTOFF-TEST( $state$ ) **then return** EVAL( $state$ )

**for each**  $s$  **in** SUCCESSORS( $state$ ) **do**

$\beta \leftarrow \text{MIN}(\beta, \text{MAX-VALUE}(s, game, \alpha, \beta))$

**if**  $\beta \leq \alpha$  **then return**  $\alpha$

**end**

**return**  $\beta$

## Deterministic games in practice

Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used an endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 443,748,401,247 positions.

Chess: Deep Blue defeated human world champion Gary Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.


Othello: human champions refuse to compete against computers, who are too good.

Go: human champions refuse to compete against computers, who are too bad. In go,  $b > 300$ , so most programs use pattern knowledge bases to suggest plausible moves.

# Nondeterministic games

E..g, in backgammon, the dice rolls determine the legal moves

Simplified example with coin-flipping instead of dice-rolling:



## Algorithm for nondeterministic games

EXPECTIMINIMAX gives perfect play

Just like MINIMAX, except we must also handle chance nodes:

...

**if** *state* is a chance node **then**

**return** average of EXPECTIMINIMAX-VALUE of SUCCESSORS(*state*)

...

A version of  $\alpha$ - $\beta$  pruning is possible

but only if the leaf values are bounded. Why??

## Nondeterministic games in practice

Dice rolls increase  $b$ : 21 possible rolls with 2 dice

Backgammon  $\approx 20$  legal moves (can be 6,000 with 1-1 roll)

$$\text{depth } 4 = 20 \times (21 \times 20)^3 \approx 1.2 \times 10^9$$

As depth increases, probability of reaching a given node shrinks

$\Rightarrow$  value of lookahead is diminished


$\alpha$ - $\beta$  pruning is much less effective

TDGAMMON uses depth-2 search + very good EVAL

$\approx$  world-champion level



## Digression: Exact values DO matter



Behaviour is preserved only by *positive linear* transformation of  $E_{VAL}$

Hence  $E_{VAL}$  should be proportional to the expected payoff

## Summary

Games are fun to work on! (and dangerous)

They illustrate several important points about AI

- ◇ perfection is unattainable  $\Rightarrow$  must approximate
- ◇ good idea to think about what to think about
- ◇ uncertainty constrains the assignment of values to states

Games are to AI as grand prix racing is to automobile design